

Problem Statement

Liver Disease Prediction by using PyCaret – Machine Learning Project

Patients with Liver disease have been continuously increasing because of excessive consumption of alcohol, inhale of harmful gases, intake of contaminated food, pickles and drugs. This dataset was used to evaluate prediction algorithms in an effort to reduce burden on doctors.

Details

This data set contains 416 liver patient records and 167 non liver patient records collected from North East of Andhra Pradesh, India. The "Dataset" column is a class label used to divide groups into liver patient (liver disease) or not (no disease). This data set contains 441 male patient records and 142 female patient records.

Any patient whose age exceeded 89 is listed as being of age "90".

Columns:

- Age of the patient
- Gender of the patient
- Total Bilirubin
- Direct Bilirubin
- Alkaline Phosphatase
- Alamine Aminotransferase
- Aspartate Aminotransferase
- Total Protiens
- Albumin
- Albumin and Globulin Ratio
- Dataset: field used to split the data into two sets (patient with liver disease, or no disease)

Objective:

The primary goal is to visualize the data and establish trends or important characteristics, if any. The next important objective is to create a pycaret model that can predict the class of the patient by the virtue of its input.

What is PyCaret and Why Should you Use it?

PyCaret is an open-source, machine learning library in Python that helps you from data preparation to model deployment.

It is easy to use and you can do almost every data science project task with just one line of code.



Task to Do:

1. Installing PyCaret on your machine and import the dataset

```
#Make necessary imports
import warnings
warnings.filterwarnings("ignore")

import pandas as pd
from pycaret.datasets import get_data

dataset = pd.read_csv("indian_liver_patient.csv")
dataset.head()
```

2. Let's start by reading the dataset using the Pandas library:

| | Age | Gender | Total_Bilirubin | Direct_Bilirubin | Alkaline_Phosphotase | Alamine_Aminotransferase | Aspartate_Aminotransferase | Total_Protiens | Albumin | Albumin_and_Globulin_Ratio | Datas |
|---|-----|--------|-----------------|------------------|----------------------|--------------------------|----------------------------|----------------|---------|----------------------------|-------|
| 0 | 65 | Female | 0.7 | 0.1 | 187 | 16 | 18 | 6.8 | 3.3 | 0.90 | |
| 1 | 62 | Male | 10.9 | 5.5 | 699 | 64 | 100 | 7.5 | 3.2 | 0.74 | |
| 2 | 62 | Male | 7.3 | 4.1 | 490 | 60 | 68 | 7.0 | 3.3 | 0.89 | |
| 3 | 58 | Male | 1.0 | 0.4 | 182 | 14 | 20 | 6.8 | 3.4 | 1.00 | |
| 4 | 72 | Male | 3.9 | 2.0 | 195 | 27 | 59 | 7.3 | 2.4 | 0.40 | |

3. The very first step before we start our machine learning project in PyCaret is to set up the environment. It's just a two-step process:

- **Importing a Module:** Depending upon the type of problem you are going to solve, you first need to import the module. In the first version of PyCaret, 6 different modules are available – regression, classification, clustering, natural language processing (NLP), anomaly detection, and associate mining rule
- **Initializing the Setup:** In this step, PyCaret performs some basic preprocessing tasks, like ignoring the IDs and Date Columns, imputing the missing values, encoding the categorical variables, and splitting the dataset into the train-test split for the rest of the modeling steps. When you run the setup function, it will first confirm the data types, and then if you press enter, it will create the environment for you to go ahead

4. Now setting environment in Pycaret by using classification module

| | Description | Value |
|----|---------------------------|------------------|
| 0 | session_id | 123 |
| 1 | Target | Dataset |
| 2 | Target Type | Binary |
| 3 | Label Encoded | 1: 0, 2: 1 |
| 4 | Original Data | (554, 11) |
| 5 | Missing Values | True |
| 6 | Numeric Features | 9 |
| 7 | Categorical Features | 1 |
| 8 | Ordinal Features | False |
| 9 | High Cardinality Features | False |
| 10 | High Cardinality Method | None |
| 11 | Transformed Train Set | (387, 10) |
| 12 | Transformed Test Set | (167, 10) |
| 13 | Shuffle Train-Test | True |
| 14 | Stratify Train-Test | False |
| 15 | Fold Generator | StratifiedKFold |
| 16 | Fold Number | 10 |
| 17 | CPU Jobs | -1 |
| 18 | Use GPU | False |
| 19 | Log Experiment | False |
| 20 | Experiment Name | clf-default-name |

PyCaret by default imputes the missing value in the dataset by 'mean' for numeric features and 'constant' for categorical features.

To change the imputation method, `numeric_imputation` and `categorical_imputation` parameters can be used within setup.

Here's a quick reminder of the evaluation metrics used for supervised learning:

- Classification: Accuracy, AUC, Recall, Precision, F1, Kappa
- Regression: MAE, MSE, RMSE, R2, RMSLE, MAPE

5. Training our machine learning model using PyCaret:

Training a model in PyCaret is quite simple. You just need to use the `create_model` function that takes just the one parameter – the model abbreviation as a string. Here, we are going to first train a LogisticRegression model

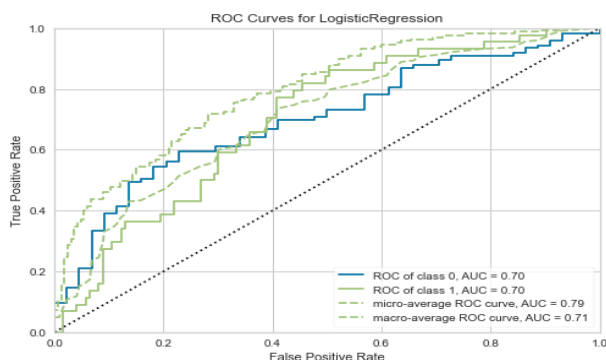
6. Hyperparameter Tuning:


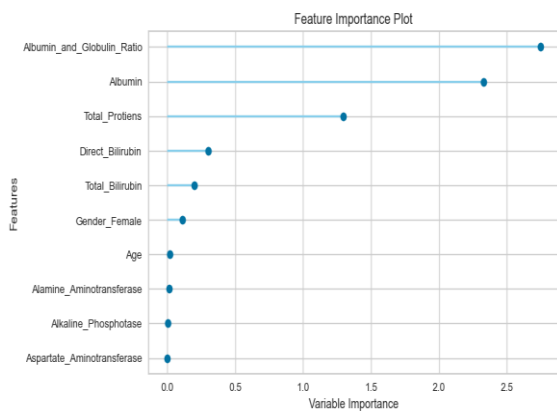
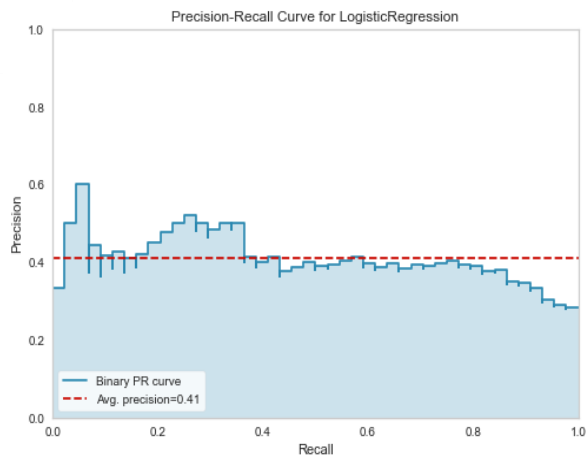
We can tune the hyperparameters of a machine learning model by just using the `tune_model` function which takes one parameter – the model abbreviation string (the same as we used in the `create_model` function).

PyCaret provides us a lot of flexibility. For example, we can define the number of folds using the `fold` parameter within the `tune_model` function. Or we can change the number of iterations using the `n_iter` parameter. Increasing the `n_iter` parameter will obviously increase the training time but will give a much better performance.

7. Plotting Metrix:

Classification: Accuracy, AUC, Recall, Precision, F1, Kappa





8. Compare model:

| | Model | Accuracy | AUC | Recall | Prec. | F1 | Kappa | MCC | TT (Sec) |
|----------|---------------------------------|----------|--------|--------|--------|--------|--------|--------|----------|
| et | Extra Trees Classifier | 0.7547 | 0.8048 | 0.4644 | 0.6744 | 0.5302 | 0.3739 | 0.3984 | 0.0290 |
| ada | Ada Boost Classifier | 0.7441 | 0.7920 | 0.4955 | 0.5697 | 0.5250 | 0.3541 | 0.3573 | 0.0130 |
| xgboost | Extreme Gradient Boosting | 0.7314 | 0.7901 | 0.4697 | 0.5622 | 0.5048 | 0.3247 | 0.3311 | 0.1620 |
| catboost | CatBoost Classifier | 0.7262 | 0.7982 | 0.4098 | 0.5976 | 0.4668 | 0.2928 | 0.3129 | 0.2830 |
| rf | Random Forest Classifier | 0.7236 | 0.7918 | 0.4273 | 0.5665 | 0.4702 | 0.2937 | 0.3062 | 0.0300 |
| lr | Logistic Regression | 0.7211 | 0.7821 | 0.3659 | 0.5817 | 0.4327 | 0.2630 | 0.2829 | 0.0140 |
| gbc | Gradient Boosting Classifier | 0.7184 | 0.7718 | 0.4091 | 0.5381 | 0.4588 | 0.2759 | 0.2831 | 0.0140 |
| lightgbm | Light Gradient Boosting Machine | 0.7183 | 0.7839 | 0.4258 | 0.5068 | 0.4571 | 0.2772 | 0.2781 | 0.0870 |
| knn | K Neighbors Classifier | 0.7134 | 0.7388 | 0.4947 | 0.5173 | 0.4985 | 0.3013 | 0.3051 | 0.0050 |
| lda | Linear Discriminant Analysis | 0.7054 | 0.7303 | 0.1644 | 0.4917 | 0.2362 | 0.1178 | 0.1423 | 0.0030 |
| ridge | Ridge Classifier | 0.7028 | 0.0000 | 0.1303 | 0.4517 | 0.1956 | 0.0921 | 0.1127 | 0.0030 |
| dt | Decision Tree Classifier | 0.6822 | 0.6313 | 0.5045 | 0.4783 | 0.4835 | 0.2577 | 0.2624 | 0.0030 |
| svm | SVM - Linear Kernel | 0.6530 | 0.0000 | 0.3212 | 0.2708 | 0.2279 | 0.1037 | 0.1282 | 0.0030 |
| nb | Naive Bayes | 0.6041 | 0.7531 | 0.9288 | 0.4249 | 0.5828 | 0.2958 | 0.3794 | 0.0030 |
| qda | Quadratic Discriminant Analysis | 0.5862 | 0.7555 | 0.9030 | 0.4114 | 0.5646 | 0.2648 | 0.3429 | 0.0040 |

ExtraTreesClassifier(bootstrap=False, ccp_alpha=0.0, class_weight=None, criterion='gini', max_depth=None, max_features='auto', max_leaf_nodes=None, max_samples=None, min_impurity_decrease=0.0, min_impurity_split=None,

9. Predict the model:

| | Model | Accuracy | AUC | Recall | Prec. | F1 | Kappa | MCC | | |
|-----|---------------------|-----------------|------------------|----------------------|--------------------------|----------------------------|----------------|---------|----------------------------|--------|
| 0 | Logistic Regression | 0.7186 | 0.6999 | 0.3636 | 0.4571 | 0.4051 | 0.2239 | 0.2264 | | |
| | Age | Total_Bilirubin | Direct_Bilirubin | Alkaline_Phosphatase | Alamine_Aminotransferase | Aspartate_Aminotransferase | Total_Protiens | Albumin | Albumin_and_Globulin_Ratio | Gender |
| 0 | 33.0 | 0.900000 | 0.800000 | 680.0 | 37.0 | 40.0 | 5.9 | 2.6 | 0.80 | |
| 1 | 45.0 | 0.800000 | 0.200000 | 140.0 | 24.0 | 20.0 | 6.3 | 3.2 | 1.00 | |
| 2 | 40.0 | 30.799999 | 18.299999 | 285.0 | 110.0 | 186.0 | 7.9 | 2.7 | 0.50 | |
| 3 | 57.0 | 0.600000 | 0.100000 | 210.0 | 51.0 | 59.0 | 5.9 | 2.7 | 0.80 | |
| 4 | 41.0 | 2.700000 | 1.300000 | 580.0 | 142.0 | 68.0 | 8.0 | 4.0 | 1.00 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 162 | 47.0 | 0.800000 | 0.200000 | 236.0 | 10.0 | 13.0 | 6.7 | 2.9 | 0.76 | |
| 163 | 39.0 | 0.600000 | 0.200000 | 188.0 | 28.0 | 43.0 | 8.1 | 3.3 | 0.60 | |
| 164 | 60.0 | 22.799999 | 12.600000 | 962.0 | 53.0 | 41.0 | 6.9 | 3.3 | 0.90 | |
| 165 | 32.0 | 32.599998 | 14.100000 | 219.0 | 95.0 | 235.0 | 5.8 | 3.1 | 1.10 | |
| 166 | 40.0 | 1.100000 | 0.300000 | 230.0 | 1630.0 | 960.0 | 4.9 | 2.8 | 1.30 | |

167 rows x 13 columns